User Guide of CKTSO-GPU

(version 20221123)

Xiaoming Chen (chenxiaoming@ict.ac.cn)

CKTSO-GPU is a GPU acceleration module of CKTSO. The package provides GPU acceleration for re-factorization and solving using CUDA. It can be called after factorization with pivoting has been performed on CPU.

CKTSO-GPU can only accelerate slightly-dense matrices but a majority of circuit matrices cannot be accelerated. For those highly-sparse matrices, just using multicore CPUs can obtain the best performance. **CKTSO-GPU is recommended when flops/NNZ(L+U) is larger than 200-300.**

1. Functions

CKTSO-GPU provides 5 functions which are easy to use. Both 32-bit and 64-bit (with `_L` in the function name) integer versions are provided. For the 32-bit integer version, only the indexes of matrix **A** use 32-bit integers, but the internal data structures for LU factors still use 64-bit integers.

(1) `CKTSO_CreateGpuAccelerator/CKTSO_L_CreateGpuAccelerator`
This function creates the CKTSO-GPU handle and also retrieves pointers to the input and output arrays. The target GPU is also set. Subsequent re-factorization and solving are performed on the selected GPU.

(2)
`CKTSO_InitializeGpuAccelerator/CKTSO_L_InitializeGpuAccelerator`
This function initializes the internal data of CKTSO-GPU, based on the internal data of CKTSO. It can be called after factorization with pivoting is performed on CPU by CKTSO (i.e., `CKTSO_Factorize/CKTSO_L_Factorize` has been called). Each time when LU factors structure has been changed (i.e., `CKTSO_Factorize/CKTSO_L_Factorize` has been called), this function must be recalled before subsequent re-factorization and solving on GPU.

(3)
`CKTSO_DestroyGpuAccelerator/CKTSO_L_DestroyGpuAccelerator`
This function frees any data associated with CKTSO-GPU and also destroys the handle.

(4) `CKTSO_GpuRefactorize/CKTSO_L_GpuRefactorize`
This function performs re-factorization on GPU. The input matrix data `ax[]` is

in host memory.

(5) `CKTSO_GpuSolve/CKTSO_L_GpuSolve`
    This function performs forward and backward substitutions on GPU. The input and output vectors `b[]` and `x[]` are both in host memory. Only the row mode is supported. **For a matrix stored in the column order, please use the transposed mode of CKTSO (i.e., setting `row0_colposi_trannega` to a negative value when calling `CKTSO_Analyze/CKTSO_L_Analyze`)**, and in this case, CKTSO-GPU will treat the matrix as row-order stored and perform transposition in every re-factorization. Only one right-side vector is supported. For multiple right-side vectors, please call this function multiple times.

2. Return Values

CKTSO-GPU uses the error codes of CKTSO. In addition, CKTSO-GPU also uses the following additional error codes.
* -51: insufficient GPU resources.
* -52: CKTSO-GPU has not been initialized.
* -53: matrix has not been re-factorized by GPU.
* If the return value is a positive integer, it means a CUDA runtime function error. For example, 2 indicates `cudaErrorMemoryAllocation`, and 35 indicates `cudaErrorInsufficientDriver`. Please refer to NVIDIA's CUDA documents to check the detailed error code.

3. Input Parameters

* `iparm[0]`: timer control. Zero disables the internal timer. A positive number enables the high-precision timer. A negative number enables the low-precision timer. Default is 0 (no timer).
* `iparm[1]`: threshold for distinguishing bulk and pipeline modes in re-factorization. Default is 128.
* `iparm[2]`: factors array allocation ratio (percentage). CKTSO-GPU allocates a bit more memory for the LU factors. If the LU factors have been changed and the size is still within the allocated length, recalling `CKTSO_InitializeGpuAccelerator/CKTSO_L_InitializeGpuAccelerator` will not reallocate the memory. Default is 110, which means 1.1.
* `iparm[3]`: whether to reallocate memories when size reduces when recalling `CKTSO_InitializeGpuAccelerator/CKTSO_L_InitializeGpuAccelerator`. Default is 0 (reusing previous memories).
* `iparm[4]`: number of blocks per multiprocessor, for re-factorization. 0 means automatic decision by CKTSO-GPU. Default is 0 (automatic decision).
* `iparm[5]`: number of threads per block, for re-factorization. 0 means automatic decision by CKTSO-GPU. Default is 0 (automatic decision).
* `iparm[6]`: number of blocks per multiprocessor, for solving. 0 means automatic

decision by CKTSO-GPU. Default is 0 (automatic decision).
* `iparm[7]`: number of threads per block, for solving. 0 means automatic decision by CKTSO-GPU. Default is 0 (automatic decision).

4. Output Parameters

* `oparm[0]`: time of `CKTSO_InitializeGpuAccelerator/CKTSO_L_InitializeGpuAccelerator`, in microsecond.
* `oparm[1]`: time of `CKTSO_GpuRefactorize/CKTSO_L_GpuRefactorize`, in microsecond.
* `oparm[2]`: time of `CKTSO_GpuSolve/CKTSO_L_GpuSolve`, in microsecond.
* `oparm[3]`: host memory usage, in bytes. It only reports the host memory usage of the CKTSO-GPU instance, excluding the CKTSO instance.
* `oparm[4]`: GPU memory usage, in bytes.
* `oparm[5]`: host memory requirement, in bytes, when -4 is returned (for the last malloc/realloc failure). It is only related to the CKTSO-GPU instance.
* `oparm[6]`: GPU memory requirement, in bytes when `cudaErrorMemoryAllocation` (2) is returned (for the last `cudaMalloc` failure).

5. System Requirements

The provided libraries are compiled on CentOS 7.9 with gcc 4.8.5 and CUDA 11.0. They require **NVIDIA driver 450.51.05 or higher**.

Supported computing capabilities of CUDA: 3.5, 3.7, 5.0, 5.2, 6.0, 6.1, 7.0, 7.5, and 8.0.

CKTSO-GPU is compatible with **CKTSO version 20221123 or higher**.

On some GPUs, it is found that L1 cache must be disabled to get correct results. Thus, libraries with L1 cache disabled are also provided. The overhead is some performance loss (about ~20%).